

Automated Fortran–C++ Bindings for Scientific Applications

Seth R Johnson

High Performance Computing Methods & Applications Team
Reactor and Nuclear Systems Division
Oak Ridge National Laboratory

ORNL is managed by UT–Battelle, LLC for the US Department of Energy

Project collaborators:
Matt Bement, Kate Evans, Tim Fuller, Andrey Prokopenko

Motivation

- C++ library developers: expand user base, more opportunities for development and follow-on funding
- Fortran scientific app developers: use newly exposed algorithms and tools for your code
- Multiphysics project integration: in-memory coupling of C++ physics code to Fortran physics code
- Transitioning application teams: bite-size migration from Fortran to C++

Background

- Exascale Computing Project: many petascale scientific app codes still are primarily Fortran
- Numerical/scientific libraries are primarily C/C++
- Expose Trilinos solver library to Fortran app developers: ForTrilinos product




SWIG: Simplified Wrapper and Interface Generator

- Generate *interfaces* to existing C and C++ code and data types so that a *target language* can invoke functions and use the data
- “Glue” code: flat C-linkage wrappers to C++ functions, corresponding interfaces in target language
- Does *not* couple target languages to other target languages
- Does *not* parse target languages or create C++ proxy wrappers

- Allegro CL
- C#
- CFFI
- CLISP
- Chicken
- D
- Go
- Guile
- Java
- Javascript
- Lua
- Modula-3
- Mzscheme
- OCAML
- Octave
- Perl
- PHP
- Python
- R
- Ruby
- Scilab
- Tcl
- UFFI


SWIG execution sequence

- SWIG reads `.i` interface file which may `%include` other C/C++ headers to parse them
 - Interface file can be as little as a couple of lines for simple interfaces
 - More difficult C++/Fortran conversions require more interface code
- SWIG generates source code files: `.cxx` and `.f90`
 - This wrapper code can be distributed like regular source files
 - Library users *don't need to know SWIG*



Control flow and data conversion

- Fortran 2003 standard defines C-compatible datatypes
- Use only Fortran-compatible ISO C datatypes
- Minimize data movement of numerical arrays



Features

- **Primitive types**
- Enumerations
- **Classes** (with inheritance)
- C strings and `std::string`
- Function pointers
- **Arrays** and `std::vector`
- **Function overloading**
- **Template instantiation**
- Compile-time constants
- **Exception handling**
- **OpenACC and Thrust**

Simple addition function: input and generated C code

```
#ifndef simplest_h
#define simplest_h
int add(int a, int b);
#endif
```

simplest.h

```
%module simplest
#include "simplest.h"
```

simplest.i

```
SWIGEXPORT int _wrap_add(
    int const *farg1,
    int const *farg2) {
    int fresult ;
    int arg1 ;
    int arg2 ;
    int result;
    arg1 = (int)(*farg1);
    arg2 = (int)(*farg2);
    result = (int)add(arg1, arg2);
    fresult = (int)(result);
    return fresult;
}
```

F/C
interface

Input
argument
conversion

Wrapped
function call

Output
argument
conversion

simplest_wrap.c

Simple addition function: generated Fortran

```
module simplest
  use, intrinsic :: ISO_C_BINDING
  implicit none
  private
  public :: add
  interface
    function swigc_add(farg1, farg2) &
      bind(C, name="_wrap_add") &
      result(fresult)
      integer(C_INT), intent(in) :: farg1
      integer(C_INT), intent(in) :: farg2
      integer(C_INT) :: fresult
    end function
  end interface
contains
  ....
```

F/C
interface

```
function swigc_add(farg1, farg2) &
  bind(C, name="_wrap_add") &
  result(fresult)
integer(C_INT), intent(in) :: farg1
integer(C_INT), intent(in) :: farg2
integer(C_INT) :: fresult
```

Output
argument
conversion

....

```
function add(a, b) &
  result(swig_result)
integer(C_INT) :: swig_result
integer(C_INT), intent(in) :: a
integer(C_INT), intent(in) :: b
integer(C_INT) :: fresult
integer(C_INT) :: farg1
integer(C_INT) :: farg2

farg1 = a
farg2 = b
fresult = swigc_add(farg1, farg2)
swig_result = fresult
end function
end module
```

Fortran
proxy
function

Input
argument
conversion

Wrapper
function
call

simplest.f90

Simple addition function: the part app devs care about

```
SWIGEXPORT int swigc_add(int const *farg1,  
int const *farg2);
```

simplest_wrap.c

```
module simplest  
  use, intrinsic :: ISO_C_BINDING  
  ...  
contains  
function add(a, b) &  
  result(swigf_result)  
  integer(C_INT) :: swigf_result  
  integer(C_INT), intent(in) :: a  
  integer(C_INT), intent(in) :: b  
  ...  
end function  
end module
```

simplest.f90

```
program main  
  use simplest, only : add  
  write (0,*) add(10,20)  
end program
```

main.f90

```
$ ./main.exe  
30
```

Templated class

```
template<typename T>
class Thing {
    T val_;
public:
    Thing(T val);
    T get() const;
};
```

```
template<typename T>
void print_thing(const Thing<T>& t);
```

Insert raw C++
code into
generated
wrapper file

```
%module "templated"
```

```
%{
#include "templated.hpp"
%}
#include "templated.hpp"
```

Tell SWIG to
parse the
header file

```
// Instantiate templated classes
%template(Thing_Int) Thing<int>;
%template(Thing_Dbl) Thing<double>;

// Instantiate and overload a function
%template(print_thing) print_thing<int>;
%template(print_thing) print_thing<double>;
```

Instantiate
templates
at SWIG
time

templated.i

Templated class: generated Fortran wrapper code

```
module templated
```

Memory ownership

```
integer, parameter, public :: swig_cmem_own_bit = 0
integer, parameter, public :: swig_cmem_rvalue_bit = 1
integer, parameter, public :: swig_cmem_const_bit = 2
type, bind(C) :: SwigClassWrapper
  type(C_PTR), public :: cptr = C_NULL_PTR
  integer(C_INT), public :: cmemflags = 0
end type
```

Opaque class wrapper

```
! class Thing< int >
type, public :: Thing_Int
  type(SwigClassWrapper), public :: swigdata
contains
  procedure :: get => swigf_Thing_Int_get
  procedure :: release => swigf_release_Thing_Int
  procedure, private :: swigf_Thing_Int_op_assign__
  generic :: assignment(=) => swigf_Thing_Int_op_assign__
end type Thing_Int
```

Fortran proxy class

```
interface Thing_Int
  module procedure swigf_create_Thing_Int
end interface
```

Second template instantiation

```
! class Thing< double >
type, public :: Thing_Dbl
...
end type Thing_Dbl
```

```
interface print_thing
  module procedure swigf_print_thing__SWIG_1,
  swigf_print_thing__SWIG_2
end interface
public :: print_thing
```

Overloaded function

```
interface
  subroutine swigc_delete_Thing_Int(farg1)
    bind(C, name="_wrap_delete_Thing_Int")
  end subroutine
end interface
```

Call delete if we "own" the memory

```
...
contains
subroutine swigf_release_Thing_Int(self)
  use, intrinsic :: ISO_C_BINDING
  class(Thing_Int), intent(inout) :: self
  type(SwigClassWrapper) :: farg1
  farg1 = self%swigdata
```

```
  if (btest(farg1%cmemflags, swig_cmem_own_bit)) then
    call swigc_delete_Thing_Int(farg1)
  endif
```

```
  farg1%cptr = C_NULL_PTR
  farg1%cmemflags = 0
  self%swigdata = farg1
```

```
end subroutine
```

```
...
end module
```

templated.f90 (2/2)

Exception handling

```
%module except
#include <std_except.i>
%exception {
    SWIG_check_unhandled_exception();
    try {
        $action
    }
    catch (const std::exception& e) {
        SWIG_exception(SWIG_RuntimeError, e.what());
    }
}
%{
#include <stdexcept>
#include <iostream>
%}
%inline {
void do_it(int i) {
    if (i < 0)
        throw std::logic_error("N00000");
    std::cout << "Yes! I got " << i
        << std::endl;
}
void do_it_again(int i) { do_it(i); }
}
```

Replaced with
wrapper call

```
program main
    use except, only : do_it, do_it_again, ierr, get_serr
    call do_it(-3)
    if (ierr /= 0) then
        write(0,*) "Got error ", ierr, ": ", get_serr()
        ierr = 0
    endif
    call do_it(2)
    call do_it(-2)
    call do_it_again(321)
end program
```

main.f90

```
Got error          -3 : In do_it(int):
N00000
Yes! I got 2
terminate called after throwing an
instance of 'std::runtime_error'
  what(): An unhandled exception
occurred before a call to
do_it_again(int); in do_it(int): N00000
```

Array views

```
%module algorithm
```

```
%{  
#include <algorithm>  
%}
```

Treat as native
Fortran array

```
%include <typemaps.i>  
%apply (SWIGTYPE *DATA, size_t SIZE) {  
    (int* data, std::size_t size) };
```

```
%inline {  
void sort(int* data, std::size_t size)  
{  
    std::sort(data, data + size);  
}  
}
```

```
subroutine sort(data)  
    use, intrinsic :: ISO_C_BINDING  
    integer(C_INT), dimension(:), target :: data  
    integer(C_INT), pointer :: farg1_view  
    type(SwigArrayWrapper) :: farg1  
    if (size(data) > 0) then  
        farg1_view => data(1)  
        farg1%data = c_loc(farg1_view)  
        farg1%size = size(data)  
    else  
        farg1%data = c_null_ptr  
        farg1%size = 0  
    end if  
    call swigc_sort(farg1)  
end subroutine
```

```
program main  
    use, intrinsic :: ISO_C_BINDING  
    use algorithm, only : sort  
    implicit none  
    integer(C_INT), dimension(6) &  
        :: test_int = (/ 3, -1, 7, 3, 1, 5 /)  
  
    call sort(test_int)  
    write (*,*) test_int  
end program
```

Thrust + OpenACC

```
%module thrustacc;

#include <openacc.i>
#include <thrust.i>

%{
#include <thrust/sort.h>
%}

inline %{
template<typename T>
void swig_thrust_sort(thrust::device_ptr<T>
DATA, size_t SIZE) {
    thrust::sort(DATA, DATA + SIZE);
}
%}

template(sort) swig_thrust_sort<int>;
template(sort) swig_thrust_sort<float>;
template(sort) swig_thrust_sort<double>;
```

```
program test_thrustacc
use thrustacc
implicit none
integer, parameter :: n = 32768
integer :: i
real, dimension(:), allocatable :: a

! Generate N uniform numbers on [0,1)
allocate(a(n))
call random_number(a)

!$acc data copy(a)
!$acc kernels
do i = 1,n
    a(i) = a(i) * 10 + i
end do
!$acc end kernels
call sort(a)
!$acc end data
write(*,*) sum(a)
end program
```

Automatic BIND(C) wrapping

```
%module bindc
```

Only generate
interface code

```
%fortranbindc;
```

```
%fortran_struct(Point)
```

Don't create
proxy class

```
%inline {  
typedef struct { float x, y, z; } Point;  
void print_point(const Point* p);  
void make_point(Point* pt,  
                const float xyz[3]);  
}
```

bindc.i

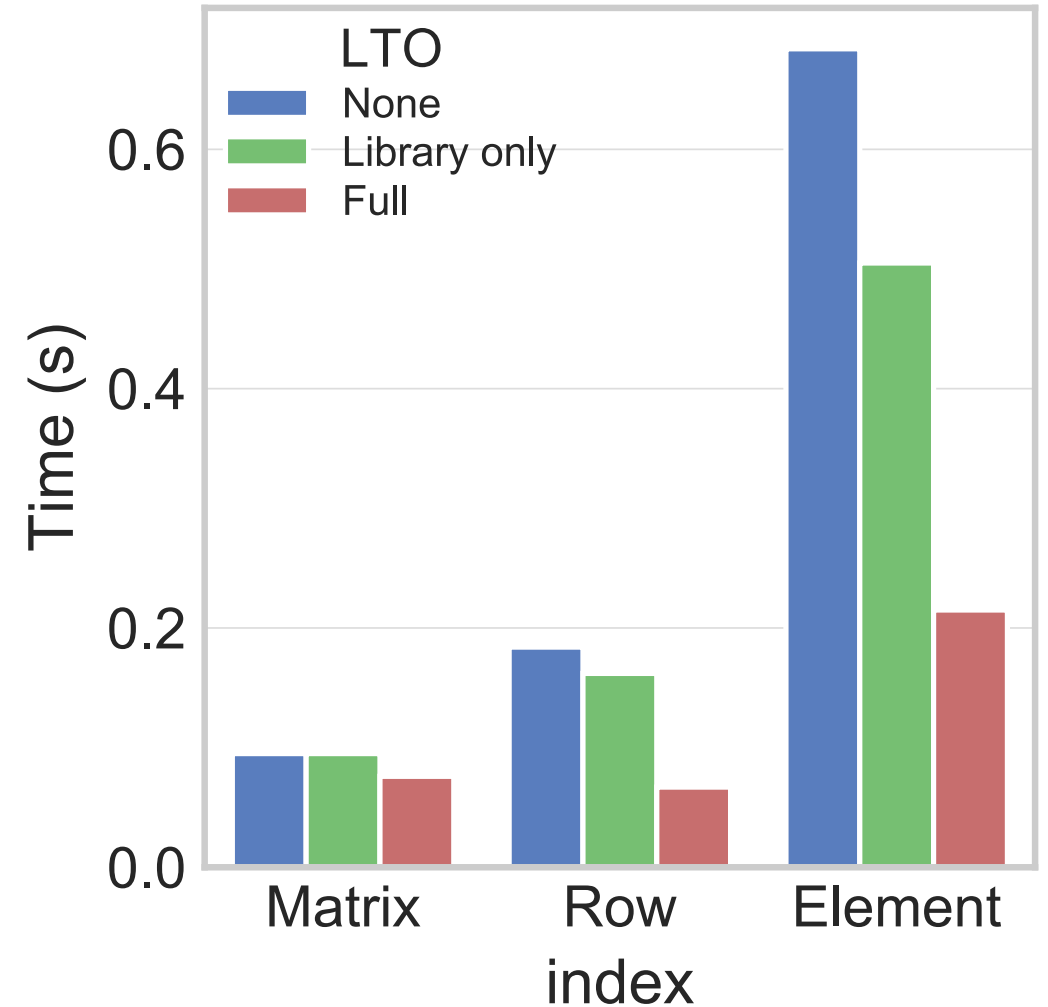
```
module bindc
```

```
...  
type, bind(C), public :: Point  
  real(C_FLOAT), public :: x  
  real(C_FLOAT), public :: y  
  real(C_FLOAT), public :: z  
end type Point  
...  
interface  
  subroutine print_point(p) &  
    bind(C, name="print_point")  
  use, intrinsic :: ISO_C_BINDING  
  import :: point  
  type(Point), intent(in) :: p  
end subroutine  
  subroutine make_point(pt, xyz) &  
    bind(C, name="make_point")  
  use, intrinsic :: ISO_C_BINDING  
  import :: point  
  type(Point) :: pt  
  real(C_FLOAT), dimension(3), intent(in) :: xyz  
end subroutine  
end interface  
end module
```

bindc.f90


Performance considerations

- Small overhead for each wrapped function call
- Link-time optimization can mitigate
- Test problem: toy numerical library with CRS matrix class
- Results: sparse matrix–vector multiply for 3000×3000 Laplacian (average of 40 runs)



Preview: Fortran bindings for C++ standard library

- Speed and reliability of C++ standard library
- Trivial installation
- App development requires only idiomatic Fortran



```
use flc_algorithm, only : argsort
implicit none
integer, dimension(5) :: iarr = [ 2, 5, -2, 3, -10000]
integer(C_INT), dimension(5) :: idx

call argsort(iarr, idx)
! This line prints a sorted array:
write(*,*) iarr(idx)
```

```
use flc_random, only : Engine, normal_distribution
real(C_DOUBLE), dimension(20) :: arr
type(Engine) :: rng

rng = Engine()
call normal_distribution(8.0d0, 2.0d0, rng, arr)
```

<https://github.com/swig-fortran/flibcpp>

Conclusions

- Implemented automatic generation of native Fortran 2003 interfaces to C++ APIs
- Supports many C++ features critical to contemporary scientific software
- Successfully generated Fortran interfaces to Trilinos numerical solvers, Sundials solvers, and C++ standard library algorithms
- High performance in typical use cases

<https://github.com/swig-fortran>

This research was supported by the Exascale Computing Project (17-SC-20-SC), a joint project of the U.S. Department of Energy's Office of Science and National Nuclear Security Administration, responsible for delivering a capable exascale ecosystem, including software, applications, and hardware technology, to support the nation's exascale computing imperative.

This research used resources of the Oak Ridge Leadership Computing Facility at the Oak Ridge National Laboratory, which is supported by the Office of Science of the U.S. Department of Energy under Contract No. DE-AC05-00OR22725.